



FEUP FACULDADE DE ENGENHARIA
UNIVERSIDADE DO PORTO

QUERY OPTIMIZATION CULTURAL FACILITIES

Database Technologies 2020/2021

Group G:

André Mori - up201700493@edu.fe.up.pt
Diogo Silva - up201405742@edu.fe.up.pt
Jiahao Zhang - up202010263@edu.fe.up.pt
Pasit Khantigul - up202010272@edu.fe.up.pt

Index

Introduction	3
Constraints	3
Indexes	4
Questions	6
Question 1	6
Query	6
Result	7
Execution Times	7
Execution Plan	8
Query using X tables	8
Comments	8
Query using Y tables	8
Comments	9
Query using Z tables	9
Comments	9
Question 2	9
Query	9
Result	10
Execution Times	10
Execution Plans	11
Query using X tables	11
Comments	11
Query using Y tables	12
Comments	12
Query using Z tables	12
Comments	13
Question 3	13
Query	13
Result	13
Execution Times	14
Query A	14
Query B	14
Execution Plan	15
Query A	15
Environment X	15
Comments	15
Environment Y	15
Comments	15
Environment Z	16

Comments	16
Query B	16
Environment X	16
Comments	17
Environment Y	17
Comments	17
Environment Z	18
Comments	18
Question 4	18
Query	18
Result	18
Execution Times	19
Execution Plan	20
Query using X tables	20
Comments	20
Query using Y tables	21
Comments	21
Query using Z tables	22
Comments	22
Question 5	22
Query	23
Result	23
Execution Times	23
Execution Plans	23
B-Tree	23
Comment	24
Bitmap	24
Comment	25
Question 6	25
Query	25
Result	25
Execution Times	26
Execution Plans	27
Query using X tables	27
Comments	27
Query using Y tables	28
Comments	28
Query using Z tables	28
Comments	28
Conclusion	29

Introduction

This project was developed in the scope of the Database Technologies course from the Master in Informatics and Computing Engineering at FEUP. The goal of the project was to develop different SQL execution plans in a database and analyse the impact of using indexes and different approaches.

To achieve the goal, three different table environments were created: X (no indexes or constraints); Y (standard integrity constraints); Z (standard integrity constraints and indexes).

Note: On the *sql/* folder delivered, there are three different files: *create.sql* to setup the environment, *indexes.sql* to create the indexes and *queries.sql* with the queries for each question.

Constraints

The constraints added to the Y and Z tables are primary and foreign keys. The code below represents the script which creates these constraints. Primary keys follow a name convention of *<TABLE_NAME>_<COLUMN>_PK*, while foreign keys follow *<TABLE_NAME>_<FOREIGN_TABLE_NAME>_FK*.

```
-- Primary Keys
ALTER TABLE "YREGIONS" ADD CONSTRAINT YREGIONS_COD_PK PRIMARY KEY("COD");
ALTER TABLE "YDISTRICTS" ADD CONSTRAINT YDISTRICTS_COD_PK PRIMARY KEY("COD");
ALTER TABLE "YMUNICIPALITIES" ADD CONSTRAINT YMUNICIPALITIES_COD_PK PRIMARY KEY("COD");
ALTER TABLE "YFACILITIES" ADD CONSTRAINT YFACILITIES_ID_PK PRIMARY KEY("ID");
ALTER TABLE "YROOMTYPES" ADD CONSTRAINT YROOMTYPES_ROOMTYPES_PK PRIMARY KEY("ROOMTYPE");
ALTER TABLE "YACTIVITIES" ADD CONSTRAINT YACTIVITIES_REF_PK PRIMARY KEY("REF");
ALTER TABLE "YUSES" ADD CONSTRAINT YUSES_ID_REF_PK PRIMARY KEY("ID","REF");
ALTER TABLE "ZREGIONS" ADD CONSTRAINT ZREGIONS_COD_PK PRIMARY KEY("COD");
ALTER TABLE "ZDISTRICTS" ADD CONSTRAINT ZDISTRICTS_COD_PK PRIMARY KEY("COD");
ALTER TABLE "ZMUNICIPALITIES" ADD CONSTRAINT ZMUNICIPALITIES_COD_PK PRIMARY KEY("COD");
ALTER TABLE "ZFACILITIES" ADD CONSTRAINT ZFACILITIES_ID_PK PRIMARY KEY("ID");
ALTER TABLE "ZROOMTYPES" ADD CONSTRAINT ZROOMTYPES_ROOMTYPES_PK PRIMARY KEY("ROOMTYPE");
ALTER TABLE "ZACTIVITIES" ADD CONSTRAINT ZACTIVITIES_REF_PK PRIMARY KEY("REF");
ALTER TABLE "ZUSES" ADD CONSTRAINT ZUSES_ID_REF_PK PRIMARY KEY("ID","REF");

-- Foreign Keys
ALTER TABLE "YDISTRICTS" ADD CONSTRAINT YDISTRICTS_REGION_FK FOREIGN KEY("REGION")
REFERENCES "YREGIONS"("COD");
ALTER TABLE "YMUNICIPALITIES" ADD CONSTRAINT YMUNICIPALITIES_REGION_FK FOREIGN
KEY("REGION") REFERENCES "YREGIONS"("COD");
ALTER TABLE "YMUNICIPALITIES" ADD CONSTRAINT YMUNICIPALITIES_DISTRICT_FK FOREIGN
KEY("DISTRICT") REFERENCES "YDISTRICTS"("COD");
ALTER TABLE "YFACILITIES" ADD CONSTRAINT YFACILITIES_MUNICIPALITY_FK FOREIGN
KEY("MUNICIPALITY") REFERENCES "YMUNICIPALITIES"("COD");
```

```

ALTER TABLE "YFACILITIES" ADD CONSTRAINT YFACILITIES_ROOMTYPE_FK FOREIGN KEY("ROOMTYPE")
REFERENCES "YROOMTYPES"("ROOMTYPE");
ALTER TABLE "YUSES" ADD CONSTRAINT YUSES_ID_FK FOREIGN KEY("ID") REFERENCES
"YFACILITIES"("ID");
ALTER TABLE "YUSES" ADD CONSTRAINT YUSES_REF_FK FOREIGN KEY("REF") REFERENCES
"YACTIVITIES"("REF");
ALTER TABLE "ZDISTRICTS" ADD CONSTRAINT ZDISTRICTS_REGION_FK FOREIGN KEY("REGION")
REFERENCES "ZREGIONS"("COD");
ALTER TABLE "ZMUNICIPALITIES" ADD CONSTRAINT ZMUNICIPALITIES_REGION_FK FOREIGN
KEY("REGION") REFERENCES "ZREGIONS"("COD");
ALTER TABLE "ZMUNICIPALITIES" ADD CONSTRAINT ZMUNICIPALITIES_DISTRICT_FK FOREIGN
KEY("DISTRICT") REFERENCES "ZDISTRICTS"("COD");
ALTER TABLE "ZFACILITIES" ADD CONSTRAINT ZFACILITIES_MUNICIPALITY_FK FOREIGN
KEY("MUNICIPALITY") REFERENCES "ZMUNICIPALITIES"("COD");
ALTER TABLE "ZFACILITIES" ADD CONSTRAINT ZFACILITIES_ROOMTYPE_FK FOREIGN KEY("ROOMTYPE")
REFERENCES "ZROOMTYPES"("ROOMTYPE");
ALTER TABLE "ZUSES" ADD CONSTRAINT ZUSES_ID_FK FOREIGN KEY("ID") REFERENCES
"ZFACILITIES"("ID");
ALTER TABLE "ZUSES" ADD CONSTRAINT ZUSES_REF_FK FOREIGN KEY("REF") REFERENCES
"ZACTIVITIES"("REF");

```

Indexes

Here we show the additional created indexes used on the Z environment to improve performance.

Query 1

```

CREATE UNIQUE INDEX INDEX_ZROOMTYPES_ROOMTYPE ON ZROOMTYPE (ROOMTYPE
ASC,DESCRIPTION ASC)
CREATE UNIQUE INDEX INDEX_ZACTIVITY_ACTIVITIES ON ZACTIVITIES (REF ASC, ACTIVITY
ASC);

```

The first index is used to access the table *ROOMTYPE* through the attribute description and project it to the attribute *room_type* and so using this index, we avoid the access to the full table.

The second index has the same functionality as the first index but it's on the table *ACTIVITIES*, since on the query 1 we need to access both the activity attribute and the ref attribute.

Query 2

```
CREATE UNIQUE INDEX INDEX_ZROOMTYPES_ROOMTYPE ON ZROOMTYPES (ROOMTYPE ASC,
DESCRIPTION ASC);
CREATE UNIQUE INDEX INDEX_ZREGIONS_DESIGNATION ON ZREGIONS (COD ASC, DESIGNATION
ASC);
CREATE UNIQUE INDEX INDEX_ZMUNICIPALITIES_REGION ON ZMUNICIPALITIES (COD ASC,
REGION ASC);
CREATE BITMAP INDEX INDEX_ZFACILITY_FACILITIES ON ZFACILITIES (MUNICIPALITY
ASC, ROOMTYPE ASC);
```

Here we use again the same index used on the 1st query on the table *ZROOMTYPE* because we need to access the table to filter the attribute description (in the WHERE condition) and we also need the *roomtype* key to do the join.

The other second and third indexes are used mostly for the join and since they involve primary keys, we can use the unique index.

The last one is the most important one because it involves 2 attributes that are not primary keys, so it's better to use a bitmap instead of a unique index, and doing the test we can see that even if the cardinality is not that low, the cost of the query is lower using the bitmap instead of b-tree index.

Query 3

```
CREATE UNIQUE INDEX INDEX_ZFACILITIES_MUNICIPALITY ON ZFACILITIES (ID,
MUNICIPALITY);
CREATE UNIQUE INDEX INDEX_ZACTIVITY_ACTIVITIES ON ZACTIVITIES (REF, ACTIVITY);
```

The indexes used for both queries 3_A and 3_B shown here are very simple : the index on *ZFACILITIES* is used mostly to avoid full access on the table for the join, and the index on *ZACTIVITIES* is used since we have a WHERE clause that involves the attribute *ACTIVITY*.

Query 4

```
CREATE UNIQUE INDEX INDEX_ZFACILITIES_MUNICIPALITY ON ZFACILITIES (ID ASC,
MUNICIPALITY ASC);
CREATE UNIQUE INDEX INDEX_ZACTIVITIES_ACTIVITY ON ZACTIVITIES (REF ASC, ACTIVITY
ASC);
CREATE UNIQUE INDEX INDEX_ZMUNICIPALITIES_DESIGNATION ON ZMUNICIPALITIES (COD
ASC, DESIGNATION ASC);
```

These 3 indexes are the same used in the last 3 queries because some access to the tables are made on attributes and not on primary keys.

Here it is not convenient to use bit-map because the attributes involved in the same table are very less and we are using the primary key as one of the attributes of the indexes, which decreases the efficiency of the bit-map indexes.

Query 5

```
---- Index for 5 a)
CREATE INDEX INDEX_B_TREE ON ZFACILITIES (ROOMTYPE, MUNICIPALITY);

---- Index for 5 b)
CREATE BITMAP INDEX INDEX_BITMAP ON ZFACILITIES (ROOMTYPE, MUNICIPALITY);
```

These indexes used for question 5 were instructed by the assignment, using *roomtype* and *municipality* columns on the *Facilities* table. For question 5 a) a B-Tree index is added and for 5 B) a Bitmap index is added. More details are presented at the Question 5 section [here](#).

Query 6

```
CREATE UNIQUE INDEX INDEX_ZMUNICIPALITIES_DISTRICT ON ZMUNICIPALITIES (COD
ASC,DISTRICT ASC);
CREATE UNIQUE INDEX INDEX_ZFACILITIES_MUNICIPALITY ON ZFACILITIES (ID ASC,
MUNICIPALITY ASC);
CREATE UNIQUE INDEX INDEX_ZDISTRICTS_COD ON ZDISTRICTS (COD ASC, DESIGNATION
ASC);
```

Most of the indexes created here are used to avoid full access to the table for the join. Here we use unique B-tree indexes and not bitmap indexes because the indexes we are creating involve primary keys and have a high cardinality, which may decrease the efficiency of the bitmap indexes.

Questions

Question 1

Which are the facilities where the room type description contains 'touros' and have 'teatro' as one of their activities? Show the id, name, description and activity.

Query

```
SELECT id, name, description, activity
```

```

FROM xfacilities NATURAL JOIN xroomtypes NATURAL JOIN xuses
NATURAL JOIN xactivities
WHERE description LIKE '%touroso%' AND activity = 'teatro';

```

Result

ID	NAME	DESCRIPTION	ACTIVITY
916	COLISEU JOSÉ RONDÃO DE ALMEIDA-EX PRAÇA DE TOIROS	Praça de touros multiusos	teatro
940	ARENA DE ÉVORA – EX PRAÇA DE TOIROS	Praça de touros multiusos	teatro
957	COLISEU DE REDONDO – EX PRAÇA DE TOIROS	Praça de touros multiusos	teatro

Execution Times

Environment	X	Y	Z
Time (s)	0.031	0.022	0.019
	0.039	0.019	0.018
	0.030	0.024	0.019
	0.055	0.031	0.019
	0.025	0.021	0.019
Average Time (s)	0.036	0.023	0.019

Execution Plan

Query using X tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			17	16
HASH JOIN			17	16
Access Predicates				
AND				
XUSES.REF=XACTIVITIES.REF				
XFACILITIES.ID=XUSES.ID				
HASH JOIN			54	12
Access Predicates				
XFACILITIES.ROOMTYPE=XROOMTYPES.ROOMTYPE				
MERGE JOIN		CARTESIAN	1	6
TABLE ACCESS	XROOMTYPES	FULL	1	3
Filter Predicates				
AND				
XROOMTYPES.DESCRPTION LIKE '%toursos%'				
XROOMTYPES.DESCRPTION IS NOT NULL				
BUFFER		SORT	1	3
TABLE ACCESS	XACTIVITIES	FULL	1	3
Filter Predicates				
XACTIVITIES.ACTIVITY='teatro'				
TABLE ACCESS	XFACILITIES	FULL	1084	6
TABLE ACCESS	XUSES	FULL	2059	4

Comments

The cost of the query for the environment X is mostly dominated by the table access on the tables *XUSES* and *XFACILITIES*, which has a lot of info. In fact we can see that the cost for the full access to the 2 tables are 10 and it's more than half of the total cost of the query.

Query using Y tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			17	12
NESTED LOOPS			17	12
HASH JOIN			54	12
Access Predicates				
YFACILITIES.ROOMTYPE=YROOMTYPES.ROOMTYPE				
MERGE JOIN		CARTESIAN	1	6
TABLE ACCESS	YROOMTYPES	FULL	1	3
Filter Predicates				
AND				
YROOMTYPES.DESCRPTION LIKE '%toursos%'				
YROOMTYPES.DESCRPTION IS NOT NULL				
BUFFER		SORT	1	3
TABLE ACCESS	YACTIVITIES	FULL	1	3
Filter Predicates				
YACTIVITIES.ACTIVITY='teatro'				
TABLE ACCESS	YFACILITIES	FULL	1084	6
INDEX	USED_ID_REF_PK	UNIQUE SCAN	1	0
Access Predicates				
AND				
YFACILITIES.ID=YUSES.ID				
YUSES.REF=YACTIVITIES.REF				

Comments

Here, the cost is significantly reduced due to the primary key index on table *YFACILITIES* and *YUSES* because the join is done with the primary keys, but we still can see that half of the cost is still dominated by the access on table *YFACILITIES*, *YACTIVITIES* and *YROOMTYPES* because we are not accessing the tables using the primary key but by their attributes

Query using Z tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			17	8
NESTED LOOPS			17	8
HASH JOIN			54	8
Access Predicates	ZFACILITIES.ROOMTYPE=ZROOMTYPES.ROOMTYPE			
MERGE JOIN		CARTESIAN	1	2
INDEX	INDEX ZROOMTYPES ROOMTYPE	FULL SCAN	1	1
Filter Predicates	ZROOMTYPES.DESCRPTION LIKE '%tours%'			
AND	ZROOMTYPES.DESCRPTION IS NOT NULL			
BUFFER		SORT	1	1
INDEX	INDEX ZACTIVITY ACTIVITIES	SKIP SCAN	1	1
Access Predicates	ZACTIVITIES.ACTIVITY='teatro'			
Filter Predicates	ZACTIVITIES.ACTIVITY='teatro'			
TABLE ACCESS	ZFACILITIES	FULL	1084	6
INDEX	ZUSES ID REF PK	UNIQUE SCAN	1	0
Access Predicates	ZFACILITIES.ID=ZUSES.ID			
AND	ZUSES.REF=ZACTIVITIES.REF			

Comments

In the environment Z, we added the 2 indexes described on the INDEX section, and we can see that the cost is reduced by $\frac{1}{3}$ by adding the 2 indexes on *ZROOMTYPES* and *ZACTIVITY*, which each one index decreases the final cost of the query by 2 values.

Question 2

How many facilities with 'tours' in the room type description are there in each region?

Query

```
SELECT xregions.designation, count(*)
FROM xfacilities NATURAL JOIN xroomtypes INNER JOIN
xmunicipalities
ON xmunicipalities.cod = xfacilities.municipality
INNER JOIN xregions ON xmunicipalities.region = xregions.cod
```

```
WHERE description LIKE '%touroso%'
GROUP BY xregions.designation;
```

Result

DESIGNATION	COUNT(*)
Lisboa	6
Norte	3
Centro	11
Alentejo	43
Algarve	1

Execution Times

Environment	X	Y	Z
Time (s)	0.031	0.022	0.020
	0.031	0.023	0.019
	0.033	0.021	0.020
	0.023	0.022	0.024
	0.025	0.023	0.019
Average Time (s)	0.028	0.022	0.020

Execution Plans

Query using X tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			7	16
HASH				
HASH JOIN		GROUP BY	7	16
Access Predicates				
XMUNICIPALITIES.REGION=XREGIONS.COD				
HASH JOIN			54	12
Access Predicates				
AND				
XMUNICIPALITIES.COD=XFACILITIES.MUNICIPALITY				
XFACILITIES.ROOMTYPE=XROOMTYPES.ROOMTYPE				
MERGE JOIN		CARTESIAN	277	6
TABLE ACCESS	XROOMTYPES	FULL	1	3
Filter Predicates				
AND				
XROOMTYPES.DESCRPTION LIKE '%touros%'				
XROOMTYPES.DESCRPTION IS NOT NULL				
BUFFER		SORT	308	3
TABLE ACCESS	XMUNICIPALITIES	FULL	308	3
TABLE ACCESS	XFACILITIES	FULL	1084	6
TABLE ACCESS	XREGIONS	FULL	7	3

Comments

In this execution plan the first operation is a full search on the table *XROOMTYPES* while filtering only tuples in which the description of the roomtypes contains 'touros'. We can see here that the cost also depends on the full accesses done within the tables: *XMUNICIPALITIES*, *XFACILITIES*, *XREGIONS*.

Query using Y tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			7	16
HASH		GROUP BY	7	16
HASH JOIN			54	15
Access Predicates	YMUNICIPALITIES.REGION=YREGIONS.COD			
NESTED LOOPS			54	15
NESTED LOOPS				
STATISTICS COLLECTOR				
HASH JOIN			54	12
Access Predicates	YMUNICIPALITIES.COD=YFACILITIES.MUNICIPALITY			
AND	YFACILITIES.ROOMTYPE=YROOMTYPES.ROOMTYPE			
MERGE JOIN		CARTESIAN	277	6
TABLE ACCESS	YROOMTYPES	FULL	1	3
Filter Predicates	YROOMTYPES.DESCRPTION LIKE '%toursos%'			
AND	YROOMTYPES.DESCRPTION IS NOT NULL			
BUFFER		SORT	308	3
TABLE ACCESS	YMUNICIPALITIES	FULL	308	3
TABLE ACCESS	YFACILITIES	FULL	1084	6
INDEX	REGION_COD_PK	UNIQUE SCAN		
Access Predicates	YMUNICIPALITIES.REGION=YREGIONS.COD			
TABLE ACCESS	YREGIONS	BY INDEX ROWID	1	3
TABLE ACCESS	YREGIONS	FULL	7	3

Comments

In this case, even by adding the primary key constraint in the table *YREGIONS*, a full search is still needed since the query is selecting the 'designation' of the regions, which is not a primary key. For what concerns the other tables, they are all fully accessed. Therefore the cost of the execution plan X and Y are the same.

- Query using Z tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			7	5
HASH		GROUP BY	7	5
NESTED LOOPS			54	4
HASH JOIN			379	4
Access Predicates	ZFACILITIES.ROOMTYPE=ZROOMTYPES.ROOMTYPE			
MERGE JOIN		CARTESIAN	6	2
INDEX	INDEX_ZROOMTYPES_ROOMTYPE	FULL SCAN	1	1
Filter Predicates	ZROOMTYPES.DESCRPTION LIKE '%toursos%'			
AND	ZROOMTYPES.DESCRPTION IS NOT NULL			
BUFFER		SORT	7	1
INDEX	INDEX_ZREGIONS_DESIGNATION	FULL SCAN	7	1
BITMAP CONVERSION		TO ROWIDS	1084	2
BITMAP INDEX	INDEX1_BITMAP	FAST FULL SCAN		
INDEX	INDEX_ZMUNICIPALITIES_REGION	UNIQUE SCAN	1	0
Access Predicates	ZMUNICIPALITIES.COD=ZFACILITIES.MUNICIPALITY			
AND	ZMUNICIPALITIES.REGION=ZREGIONS.COD			

Comments

In the last environment, we can see that the cost has been significantly decreased due to the indexes used (shown in the indexes paragraph). The reason why one of the indexes is a bit-map is because the attributes where we apply the bit-map indexes are not primary keys. Even if the cardinality is not as shown at lecture less than 1/100, we saw that it still has a good effort on decreasing the cost of the query.

Question 3

How many municipalities do not have any facility with an activity of 'cinema'?

- a. Use not in.*
- b. Use external join and is null.*

Query

```
---- a)
SELECT count(*) FROM
xmunicipalities
WHERE cod NOT IN
(SELECT municipality
FROM xfacilities NATURAL JOIN xuses NATURAL JOIN xactivities
INNER JOIN
xmunicipalities ON xmunicipalities.cod = xfacilities.municipality
WHERE activity = 'cinema');

---- b)
SELECT count(*)
FROM xmunicipalities LEFT OUTER JOIN (SELECT municipality as cod,
activity
FROM xfacilities NATURAL JOIN xuses NATURAL JOIN xactivities
INNER JOIN
xmunicipalities ON xmunicipalities.cod = xfacilities.municipality
WHERE activity = 'cinema') cinemas
ON xmunicipalities.cod = cinemas.cod
WHERE activity IS NULL;
```

Result

COUNT(*)
100

Execution Times

Query A

Environment	X	Y	Z
Time (s)	0.028	0.023	0.023
	0.026	0.022	0.023
	0.028	0.024	0.023
	0.025	0.027	0.021
	0.027	0.021	0.020
Average Time (s)	0.027	0.023	0.022

Query B

Environment	X	Y	Z
Time (s)	0.024	0.022	0.023
	0.028	0.022	0.023
	0.023	0.024	0.021
	0.027	0.022	0.021
	0.025	0.022	0.023
Average Time (s)	0.025	0.022	0.022

Execution Plan

Query A

Environment X

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				19
SORT		AGGREGATE		1
HASH JOIN		ANTI SNA		197
Access Predicates				
COD=MUNICIPALITY				
TABLE ACCESS	XMUNICIPALITIES	FULL		308
VIEW	SYS.VIW_NSO_1			343
HASH JOIN		RIGHT SEMI		343
Access Predicates				
XMUNICIPALITIES.COD=XFACILITIES.MUNICIPALITY				
TABLE ACCESS	XMUNICIPALITIES	FULL		308
HASH JOIN		SEMI		343
Access Predicates				
AND				
XUSES.REF=XACTIVITIES.REF				
XFACILITIES.ID=XUSES.ID				
MERGE JOIN		CARTESIAN		1084
TABLE ACCESS	XACTIVITIES	FULL		1
Filter Predicates				
XACTIVITIES.ACTIVITY='cinema'				
BUFFER		SORT		1084
TABLE ACCESS	XFACILITIES	FULL		1084
TABLE ACCESS	XUSES	FULL		2059

Comments

In this environment the high cost is due to the fact that a full search is used 5 times (including twice on the table *MUNICIPALITIES*).

Environment Y

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				10
SORT		AGGREGATE		1
HASH JOIN		ANTI		197
Access Predicates				
COD=MUNICIPALITY				
INDEX	YMUNICIPALITIES_COD_PK	FULL SCAN		308
VIEW	SYS.VIW_NSO_1			343
NESTED LOOPS		SEMI		343
MERGE JOIN		CARTESIAN		1084
TABLE ACCESS	YACTIVITIES	FULL		1
Filter Predicates				
YACTIVITIES.ACTIVITY='cinema'				
BUFFER		SORT		1084
TABLE ACCESS	YFACILITIES	FULL		1084
Filter Predicates				
YFACILITIES.MUNICIPALITY IS NOT NULL				
INDEX	YUSES_ID_REF_PK	UNIQUE SCAN		2055
Access Predicates				
AND				
YFACILITIES.ID=YUSES.ID				
YUSES.REF=YACTIVITIES.REF				

Comments

The biggest change here is the number of full searches that is reduced in more than half. The search on the *USES* table turns into an unique scan, due to the fact that we now have primary and foreign keys, and there is no longer needed two full searches on the table *MUNICIPALITIES*. Instead only a full scan is used.

Environment Z

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	5
SORT		AGGREGATE	1	1
HASH JOIN		ANTI	197	5
Access Predicates				
COD=MUNICIPALITY				
INDEX	ZMUNICIPALITIES_COD_PK	FULL SCAN	308	1
VIEW	SYS.V_W_NSQ_1		343	4
NESTED LOOPS		SEMI	343	4
MERGE JOIN		CARTESIAN	1084	4
INDEX	INDEX_ZACTIVITY_ACTIVITIES	SKIP SCAN	1	1
Access Predicates				
ZACTIVITIES.ACTIVITY='cinema'				
Filter Predicates				
ZACTIVITIES.ACTIVITY='cinema'				
BUFFER		SORT	1084	3
INDEX	INDEX_ZFACILITIES_MUNICIPALITY	FAST FULL SCAN	1084	3
Filter Predicates				
ZFACILITIES.MUNICIPALITY IS NOT NULL				
INDEX	ZUSES_ID_REF_PK	UNIQUE SCAN	2055	0
Access Predicates				
AND				
ZFACILITIES.ID=ZUSES.ID				
ZUSES.REF=ZACTIVITIES.REF				

Comments

With the introduction of the two indexes we now have reduced the cost by half. The index on the *FACILITIES* table makes it so that a full search is no longer needed and only a fast full scan is performed. The index on the *ACTIVITIES* table allows for a skip scan to take place instead of a full search.

Query B

Environment X

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	19
SORT		AGGREGATE	1	1
FILTER				
Filter Predicates				
CINEMAS.ACTIVITY IS NULL				
HASH JOIN		OUTER	1	19
Access Predicates				
XMUNICIPALITIES.COD=CINEMAS.COD(+)				
TABLE ACCESS	XMUNICIPALITIES	FULL	308	3
VIEW			343	16
HASH JOIN			343	16
Access Predicates				
XMUNICIPALITIES.COD=XFACILITIES.MUNICIPALITY				
TABLE ACCESS	XMUNICIPALITIES	FULL	308	3
HASH JOIN			343	13
Access Predicates				
AND				
XUSES.REF=XACTIVITIES.REF				
XFACILITIES.ID=XUSES.ID				
MERGE JOIN		CARTESIAN	1084	9
TABLE ACCESS	XACTIVITIES	FULL	1	3
Filter Predicates				
XACTIVITIES.ACTIVITY='cinema'				
BUFFER		SORT	1084	6
TABLE ACCESS	XFACILITIES	FULL	1084	6
TABLE ACCESS	XUSES	FULL	2059	4

Comments

In the environment X, the cost is very high and we can see that this is mostly due to the amount of full access that are performed to the tables, especially to the table *XFACILITIES* and *XUSES* (just these two use up 2/3 of the total cost).

Environment Y

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	10
SORT		AGGREGATE	1	
FILTER				
Filter Predicates				
CINEMAS.ACTIVITY IS NULL				
HASH JOIN		OUTER	1	10
Access Predicates				
YMUNICIPALITIES.COD=CINEMAS.COD(+)				
INDEX	MUNICIPALITIES_COD_PK	FULL SCAN	308	1
VIEW			343	9
NESTED LOOPS			343	9
MERGE JOIN		CARTESIAN	1084	9
TABLE ACCESS	YACTIVITIES	FULL	1	3
Filter Predicates				
YACTIVITIES.ACTIVITY='cinema'				
BUFFER		SORT	1084	6
TABLE ACCESS	YFACILITIES	FULL	1084	6
Filter Predicates				
YFACILITIES.MUNICIPALITY IS NOT NULL				
INDEX	USED_ID_REF_PK	UNIQUE SCAN	1	0
Access Predicates				
AND				
YFACILITIES.ID=YUSES.ID				
YUSES.REF=YACTIVITIES.REF				

Comments

Here the cost is decreased because we are using indexes on the primary keys on table *YUSES*, *YMUNICIPALITIES* but we can still notice that there are 2 tables with the full access that can be optimized, especially the table *YFACILITIES*, where one access has a cost of 6, which is more than the half of the total cost.

Environment Z

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	5
SORT		AGGREGATE	1	
FILTER				
Filter Predicates				
CINEMAS.ACTIVITY IS NULL				
HASH JOIN		OUTER	1	5
Access Predicates				
ZMUNICIPALITIES.COD=CINEMAS.COD(+)				
INDEX	ZMUNICIPALITIES_COD_PK	FULL SCAN	308	1
VIEW			343	4
NESTED LOOPS			343	4
MERGE JOIN		CARTESIAN	1084	4
INDEX	INDEX_ZACTIVITY_ACTIVITIES	SKIP SCAN	1	1
Access Predicates				
ZACTIVITIES.ACTIVITY='cinema'				
Filter Predicates				
ZACTIVITIES.ACTIVITY='cinema'				
BUFFER		SORT	1084	3
INDEX	INDEX_ZFACILITIES_MUNICIPAL...	FAST FULL SCAN	1084	3
Filter Predicates				
ZFACILITIES.MUNICIPALITY IS NOT NULL				
INDEX	ZUSES_ID_REF_PK	UNIQUE SCAN	1	0
Access Predicates				
AND				
ZFACILITIES.ID=ZUSES.ID				
ZUSES.REF=ZACTIVITIES.REF				

Comments

In this last environment, we added an index on the *ZFACILITIES* and another on the *ZACTIVITIES*, where both decreased almost by half their cost to access the table (comparing the execution to the previous Y environment).

Question 4

Which is the municipality with more facilities engaged in each of the six kinds of activities? Show the activity, the municipality name and the corresponding number of facilities.

Query

```
SELECT q2.designation, q1.activity, q1.facilities
FROM (SELECT activity, max(facilities) as facilities
FROM (SELECT municipality, activity, count(*) as facilities
FROM xfactivities NATURAL JOIN xuses NATURAL JOIN xactivities
GROUP BY municipality, activity)
GROUP BY activity) q1 LEFT JOIN (SELECT designation, activity,
count(*) as facilities
FROM xmunicipalities INNER JOIN xfactivities NATURAL JOIN xuses
```

```

NATURAL JOIN xactivities
ON xmunicipalities.cod = xfacilities.municipality
GROUP BY designation, activity) q2
ON q2.activity = q1.activity AND q2.facilities = q1.facilities;

```

Result

DESIGNATION	ACTIVITY	FACILITIES
Lisboa	circo	2
Lisboa	dança	47
Moura	tauromaquia	4
Lisboa	música	77
Lisboa	teatro	66
Lisboa	cinema	96

Execution Times

Environment	X	Y	Z
Time (s)	0.029	0.031	0.028
	0.032	0.031	0.029
	0.030	0.027	0.025
	0.025	0.027	0.029
	0.034	0.030	0.026
Average Time	0.030	0.029	0.0274

Execution Plan

Query using X tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			13	31
HASH JOIN		OUTER	13	31
Access Predicates				
AND				
Q2.FACILITIES(+)=Q1.FACILITIES				
Q2.ACTIVITY(+)=Q1.ACTIVITY				
VIEW			6	14
HASH		GROUP BY	6	14
VIEW			964	14
HASH		GROUP BY	964	14
HASH JOIN			2059	13
Access Predicates				
XFACILITIES.ID=XUSES.ID				
TABLE ACCESS	XFACILITIES	FULL	1084	6
HASH JOIN			2059	7
Access Predicates				
XUSES.REF=XACTIVITIES.REF				
TABLE ACCESS	XACTIVITIES	FULL	6	3
TABLE ACCESS	XUSES	FULL	2059	4
VIEW			1299	17
HASH		GROUP BY	1299	17
HASH JOIN			2059	16
Access Predicates				
XUSES.REF=XACTIVITIES.REF				
TABLE ACCESS	XACTIVITIES	FULL	6	3
HASH JOIN			2059	13
Access Predicates				
XFACILITIES.ID=XUSES.ID				
HASH JOIN			1084	9
Access Predicates				
XMUNICIPALITIES.COD=XFACILITIES.MUNICIPALITY				
TABLE ACCESS	XMUNICIPALITIES	FULL	308	3
TABLE ACCESS	XFACILITIES	FULL	1084	6
TABLE ACCESS	XUSES	FULL	2059	4

Comments

In this case, the query requested is a little bit tricky, so the solution proposed uses 3 views as support. Since 2 of the 3 views do very similar things, the cost is inevitably increased, but it's also noticeable that with some indexes the cost of table access can be decreased, and so the total cost of the query.

Query using Y tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			13	29
HASH JOIN		OUTER	13	29
Access Predicates				
AND				
	Q2.FACILITIES(+)=Q1.FACILITIES			
	Q2.ACTIVITY(+)=Q1.ACTIVITY			
VIEW			6	13
HASH		GROUP BY	6	13
VIEW			964	13
HASH		GROUP BY	964	13
HASH JOIN			2059	12
Access Predicates				
	YFACILITIES.ID=YUSES.ID			
TABLE ACCESS	YFACILITIES	FULL	1084	6
HASH JOIN			2059	6
Access Predicates				
	YUSES.REF=YACTIVITIES.REF			
TABLE ACCESS	YACTIVITIES	FULL	6	3
INDEX	USED_ID_REF_PK	FAST FULL SCAN	2059	3
VIEW			1299	16
HASH		GROUP BY	1299	16
HASH JOIN			2059	15
Access Predicates				
	YUSES.REF=YACTIVITIES.REF			
TABLE ACCESS	YACTIVITIES	FULL	6	3
HASH JOIN			2059	12
Access Predicates				
	YFACILITIES.ID=YUSES.ID			
NESTED LOOPS			2059	12
STATISTICS COLLECTOR				
HASH JOIN			1084	9
Access Predicates				
	YMUNICIPALITIES.COD=YFACILITIES.MUNICIPALITY			
TABLE ACCESS	YMUNICIPALITIES	FULL	308	3
TABLE ACCESS	YFACILITIES	FULL	1084	6
INDEX	USED_ID_REF_PK	RANGE SCAN	2	3
Access Predicates				
	YFACILITIES.ID=YUSES.ID			
INDEX	USED_ID_REF_PK	FAST FULL SCAN	2059	3

Comments

In the environment Y, we can see that the cost is a bit less with the introduction of indexes on the primary key like on table *YUSED*, but we see that there are still a lot of full access to the tables like on *YFACILITIES* and *YMUNICIPALITIES* that can be optimized by using appropriate indexes as shown in the environment Z

Query using Z tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			13	18
HASH JOIN		OUTER	13	18
Access Predicates				
AND				
Q2.FACILITIES(+)=Q1.FACILITIES				
Q2.ACTIVITY(+)=Q1.ACTIVITY				
VIEW			6	8
HASH		GROUP BY	6	8
VIEW			964	8
HASH		GROUP BY	964	8
HASH JOIN			2059	7
Access Predicates				
ZFACILITIES.ID=ZUSES.ID				
INDEX	INDEX_ZFACILITIES_MUNICIPAL...	FAST FULL SCAN	1084	3
HASH JOIN			2059	4
Access Predicates				
ZUSES.REF=ZACTIVITIES.REF				
INDEX	INDEX_ZACTIVITIES_ACTIVITY	FULL SCAN	6	1
INDEX	ZUSES_ID_REF_PK	FAST FULL SCAN	2059	3
VIEW			1299	10
HASH		GROUP BY	1299	10
HASH JOIN			2059	9
Access Predicates				
ZUSES.REF=ZACTIVITIES.REF				
INDEX	INDEX_ZACTIVITIES_ACTIVITY	FULL SCAN	6	1
HASH JOIN			2059	8
Access Predicates				
ZFACILITIES.ID=ZUSES.ID				
NESTED LOOPS			2059	8
STATISTICS COLLEC				
HASH JOIN			1084	5
Access Predicates				
ZMUNICIPALITIES.COD=ZFACILITIES.MUNICIPALITY				
INDEX	INDEX_ZMUNICIPALITIES_DESIG...	FAST FULL SCAN	308	2
INDEX	INDEX_ZFACILITIES_MUNICIPAL...	FAST FULL SCAN	1084	3
INDEX	ZUSES_ID_REF_PK	RANGE SCAN	2	3
Access Predicates				
ZFACILITIES.ID=ZUSES.ID				
INDEX	ZUSES_ID_REF_PK	FAST FULL SCAN	2059	3

Comments

Here we can see that the cost has been decreased by $\frac{2}{3}$ comparing it to the cost of the environment Y, and that could be done only by introducing indexes on table *ZMUNICIPALITIES* and *ZFACILITIES*.

The total cost, compared to the other queries, is still very high since the query is a bit tricky but reducing almost to half the cost from environment X to environment Z using adequate indexes is a very good result.

Question 5

Compare the execution plans (just the Z environment) for the query giving the municipality designation, the facility name, and the room type description, where the

room type of the facility includes 'touros' and the municipality is part of the Porto district.

a. With a B-tree index on the room type and municipality columns of the facilities table;

b. With a bitmap index on the room type and municipality columns of the facilities table.

Query

```
SELECT zmunicipalities.designation, zfacilities.name,
zroomtypes.description
FROM zfacilities NATURAL JOIN zroomtypes INNER JOIN
zmunicipalities On
zmunicipalities.cod = zfacilities.municipality INNER JOIN
zdistricts ON
zdistricts.cod = zmunicipalities.district
WHERE zdistricts.designation = 'Porto' AND zroomtypes.description
LIKE '%touros%';
```

Result

DESIGNATION	NAME	DESCRIPTION
Póvoa de Varzim	PRAÇA DE TOIROS DA PÓVOA DE VARZIM	Praça de touros

Execution Times

Index	B-TREE	BITMAP
Time (s)	0.026	0.024
	0.024	0.028
	0.025	0.025
	0.026	0.027
	0.023	0.029
Average Time	0.025	0.027

Execution Plans

The main difference between both execution plans below is that, even though both indexes use the same columns/table and have the same cost, the Bitmap index is

not even used. That must be because of the cardinality of the *Facilities* table which is very high, making the B-Tree more effective.

B-Tree

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			3	15
HASH JOIN			3	15
Access Predicates				
AND				
ZMUNICIPALITIES.COD=ZFACILITIES.MUNICIPALITY				
ZFACILITIES.ROOMTYPE=ZROOMTYPES.ROOMTYPE				
NESTED LOOPS			3	15
NESTED LOOPS				
STATISTICS COLLECTOR				
HASH JOIN			14	9
Access Predicates				
ZDISTRICTS.COD=ZMUNICIPALITIES.DISTRICT				
TABLE ACCESS ZDISTRICTS	ZDISTRICTS	FULL	1	3
Filter Predicates				
ZDISTRICTS.DESIGNATION='Porto'				
MERGE JOIN		CARTESIAN	277	6
TABLE ACCESS ZROOMTYPES	ZROOMTYPES	FULL	1	3
Filter Predicates				
AND				
ZROOMTYPES.DESCRPTION LIKE '%toursos%'				
ZROOMTYPES.DESCRPTION IS NOT NULL				
BUFFER		SORT	308	3
TABLE ACCESS ZMUNICIPALITIES	ZMUNICIPALITIES	FULL	308	3
INDEX	INDEX B TREE	RANGE SCAN		
Access Predicates				
AND				
ZFACILITIES.ROOMTYPE=ZROOMTYPES.ROOMTYPE				
ZMUNICIPALITIES.COD=ZFACILITIES.MUNICIPALITY				
TABLE ACCESS ZFACILITIES	ZFACILITIES	BY INDEX ROWID	1	6
TABLE ACCESS ZFACILITIES	ZFACILITIES	FULL	1084	6

Comment

As we can see, the access to the *ROOMTYPE* and *MUNICIPALITY* on the join clause is done using the index. The only problem here is that it is doing it twice probably because some of the attributes are not part of the B-tree index.

Bitmap

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			3	15
HASH JOIN			3	15
Access Predicates				
AND				
ZMUNICIPALITIES.COD=ZFACILITIES.MUNICIPALITY				
ZFACILITIES.ROOMTYPE=ZROOMTYPES.ROOMTYPE				
HASH JOIN			14	9
Access Predicates				
ZDISTRICTS.COD=ZMUNICIPALITIES.DISTRICT				
TABLE ACCESS	ZDISTRICTS	FULL	1	3
Filter Predicates				
ZDISTRICTS.DESIGNATION='Porto'				
MERGE JOIN		CARTESIAN	277	6
TABLE ACCESS	ZROOMTYPES	FULL	1	3
Filter Predicates				
AND				
ZROOMTYPES.DESCRPTION LIKE '%toursos%'				
ZROOMTYPES.DESCRPTION IS NOT NULL				
BUFFER		SORT	308	3
TABLE ACCESS	ZMUNICIPALITIES	FULL	308	3
TABLE ACCESS	ZFACILITIES	FULL	1084	6

Comment

Using the bitmap index we can see that it is not used and the cost remains the same.

Question 6

Which are the codes and designations of the districts with facilities in all the municipalities?

Query

```
SELECT cod, designation FROM xdistricts WHERE cod NOT IN
(
SELECT xdistricts.cod
FROM xmunicipalities INNER JOIN xdistricts ON
xmunicipalities.district = xdistricts.cod
LEFT OUTER JOIN xfacilities on xfacilities.municipality =
xmunicipalities.cod
WHERE id is null
);
```

Result

COD	DESIGNATION
15	Setúbal
7	Évora
11	Lisboa
12	Portalegre

Execution Times

Environment	X	Y	Z
Time (s)	0.026	0.022	0.019
	0.034	0.021	0.020
	0.035	0.020	0.018
	0.031	0.022	0.020
	0.028	0.035	0.019
Average Time	0.031	0.024	0.019

Execution Plans

Query using X tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			19	15
HASH JOIN		ANTI SNA	19	15
Access Predicates				
COD=COD				
TABLE ACCESS	<u>XDISTRICKTS</u>	FULL	20	3
VIEW	<u>SYS.VW_NSO_1</u>		1	12
HASH JOIN			1	12
Access Predicates				
XMUNICIPALITIES.DISTRICT=XDISTRICKTS.COD				
FILTER				
Filter Predicates				
XFACILITIES.ID IS NULL				
HASH JOIN		OUTER	1	9
Access Predicates				
XFACILITIES.MUNICIPALITY(+)=XMUNICIPALITIES.COD				
TABLE ACCESS	<u>XMUNICIPALITIES</u>	FULL	308	3
TABLE ACCESS	<u>XFACILITIES</u>	FULL	1084	6
TABLE ACCESS	<u>XDISTRICKTS</u>	FULL	20	3

Comments

As in previous queries, in the environment X, tables are fully accessed. First, a full search of *XMUNICIPALITIES* and *XFACILITIES* are done. Then only the district codes of tuples containing *null* values in the *Facilities' ID* field are selected. Using *NOT IN* operator, these results are the ones that are not considered during the full access on the table *XDISTRICKTS*.

Query using Y tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	12
MERGE JOIN		ANTI	1	12
TABLE ACCESS	YDISTRICTS	BY INDEX ROWID	20	2
INDEX	DISTRICTS_COD_PK	FULL SCAN	20	1
SORT		UNIQUE	197	10
Access Predicates				
COD=COD				
Filter Predicates				
COD=COD				
VIEW	SYS.VW_NSO_1		197	9
FILTER				
Filter Predicates				
YFACILITIES.ID IS NULL				
HASH JOIN		OUTER	197	9
Access Predicates				
YFACILITIES.MUNICIPALITY(+)=YMUNICIPALITIES.COD				
TABLE ACCESS	YMUNICIPALITIES	FULL	308	3
Filter Predicates				
YMUNICIPALITIES.DISTRICT IS NOT NULL				
TABLE ACCESS	YFACILITIES	FULL	1084	6

Comments

In the environment Y for this query the cost is reduced due to the introduction of the primary key in the *YDISTRICTS* table, named 'cod'. The reduction of the cost is not significant, since full accesses are still done on the tables *YMUNICIPALITIES* and *YFACILITIES*. To improve the search on these tables we should add specific indexes.

Query using Z tables

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	5
HASH JOIN		ANTI	1	5
Access Predicates				
COD=COD				
INDEX	INDEX_ZDISTRICTS_COD	FULL SCAN	20	1
VIEW	SYS.VW_NSO_1		197	4
FILTER				
Filter Predicates				
ZFACILITIES.ID IS NULL				
HASH JOIN		OUTER	197	4
Access Predicates				
ZFACILITIES.MUNICIPALITY(+)=ZMUNICIPALITIES.COD				
INDEX	INDEX_ZMUNICIPALITIES_DISTR...	FULL SCAN	308	1
Filter Predicates				
ZMUNICIPALITIES.DISTRICT IS NOT NULL				
INDEX	INDEX_ZFACILITIES_MUNICIPAL...	FAST FULL SCAN	1084	3

Comments

In the last environment the cost is definitely improved since it is reduced by $\frac{1}{3}$ compared to the cost from environment X and almost $\frac{1}{3}$ compared to the cost from the environment Y. From this, it was our understanding that the cost is reduced especially when appropriate indexes are introduced, in this case, on the table *ZFACILITIES* and *ZMUNICIPALITIS*.

Conclusion

With the development of the project, the group learned more about the impact of indexes on queries and how it can benefit the overall performance. Still, the indexes must be carefully selected, especially regarding its type (B-Tree or Bitmap), since it may not be effective in some circumstances and just create additional costs.